
Indexing Algorithms and Data Structures

David Boyuka, daboyuka@ncsu.edu

North Carolina State University

March 31, 2015

What Is an Index?

- A **data structure** to **speed up** certain types of **access** to a dataset vs. a **sequential scan**
 - **Minimizes costly access** to **slow storage** (disk, RAM?)
- Data consists of “**records**” (sometimes “**documents**”)
 - Generally, each record has an ID, known as an “**RID**”
- Data has “**attributes**” (“**columns**,” “**variables**,” “**fields**”)
 - Each record may have a value for each attribute
 - Attributes in **semi-structured** data **sparsely cover records**

Where Are Indexes* Used?

- Database Management Systems (DBMSs)
 - Relational
 - Distributed
 - Key-value store
- Web and Document Search
- File Systems

**While the plural “indices” of index is more common in general usage, “indexes” is common in certain contexts, including database literature.*

Overview

- **“Traditional” DBMS indexes**
- **Bitmap indexes**
- **Inverted indexes**


Overview

- **“Traditional” DBMS indexes**
- Bitmap indexes
- Inverted indexes

“Traditional” DBMS Indexes


- **B+ Tree Index**

- Based on the B-tree search tree
- Fast range and equality queries



- **Hash Index**

- Based on the hash table
- Very fast equality queries

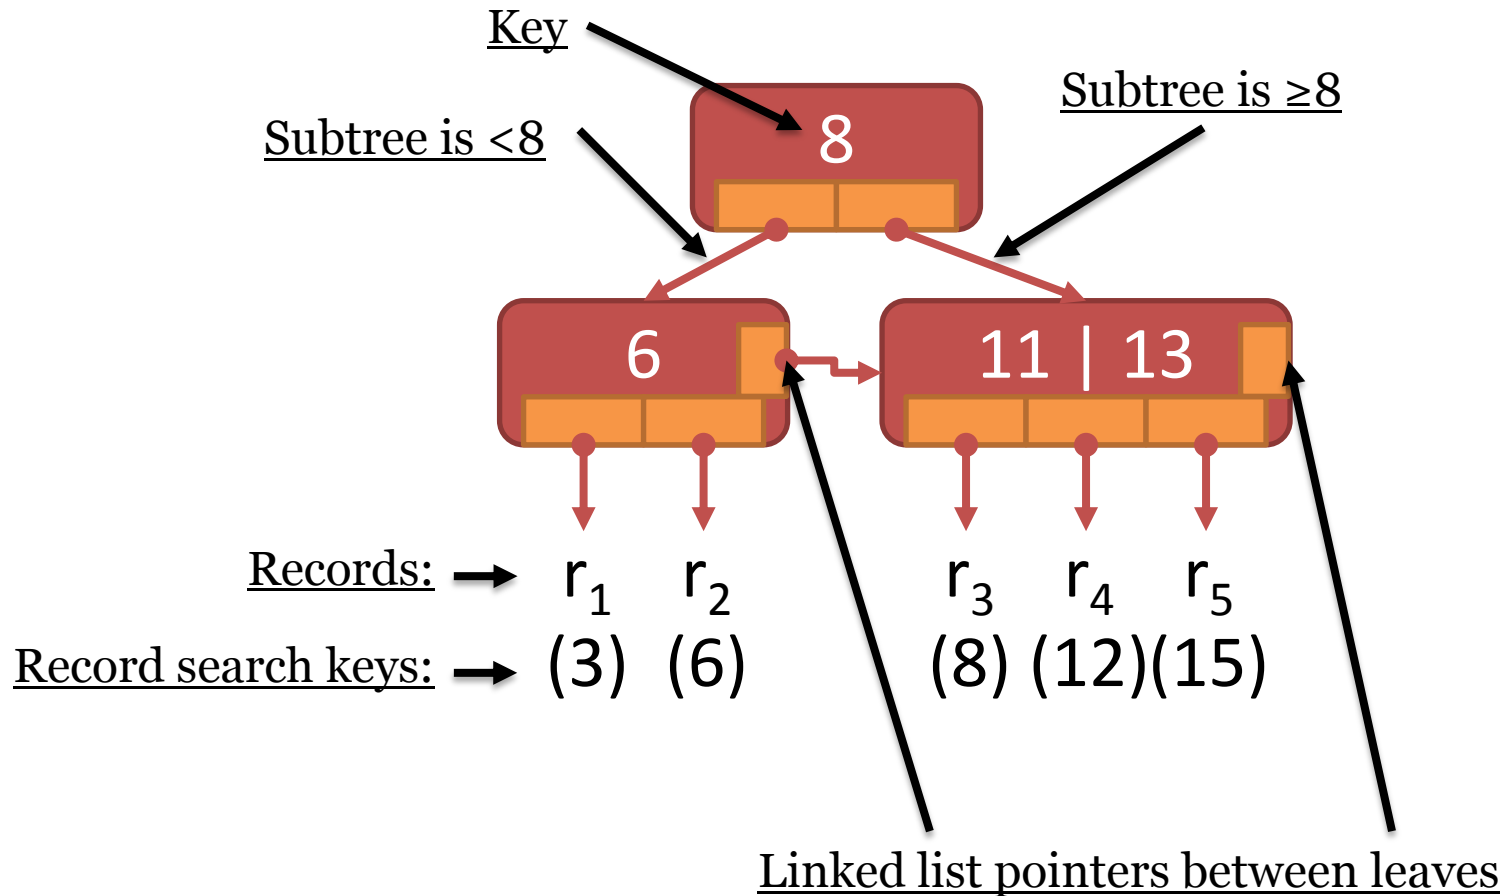


- Both map **search keys** to **matching record(s)**

The B+ Tree

- **Perfectly-balanced** search tree with a **wide fan-out**
 - Allows search over some **record attribute** (“search key”)
 - “**Order**” = b adjustable, each node has $\text{ceil}(b/2)$ to b keys
- Leaf nodes hold k **search keys**, $k+1$ **records**
 - $(\text{record } i) < (\text{key } i) \leq (\text{record } i+1)$
 - Leaf nodes connected via linked list
- Internal nodes have k **search keys**, $k+1$ **child pointers**
 - $(\text{subtree } i) < (\text{key } i) \leq (\text{subtree } i+1)$

B+ Tree Example




Inserting a Record into a B+ Tree

1. Find the leaf where the record belongs
2. Add the record and corresponding search key
3. If node has $\leq b$ records, done!
4. Else:
 - a. Split the node, add the new half to the parent node
 - b. Push the least key of the greater half to the parent
 - c. If the parent has $\leq b$ records, done!
 - d. Else, repeat from step 4a for the parent
 - e. If the root splits, add a new root with two children


B+ Tree Insert Example

Insert r_6 with search key 14




B+ Tree Insert Example

Insert r_6 with search key 14




B+ Tree Insert Example

Insert r_6 with search key 14



B+ Tree Insert Example

Insert r_6 with search key 14



Deleting a Record from a B+ Tree

1. Find the leaf containing the record
2. Delete the record and the corresponding key
3. If node has $\geq \text{ceil}(b/2)$ records, done!
4. Else:
 - a. Try stealing a key/record from an adjacent sibling
 - b. If too few keys in all siblings, merge with a sibling
 - c. If merged, repeat from 3 for parent node

(We'll skip the example and move on to other topics)

Clustered vs. Non-clustered B+ Tree


- **Non-clustered:**
 - Only **pointers to records** are stored, with full data stored elsewhere on disk
- **Clustered:**
 - Full records are stored directly in the leaves
- Maximum of **one** clustered index **per table/dataset**

B+ Trees: Pros and Cons

- + Supports both **range** and **equality queries**
- + **Inserts/deletes inexpensive** (usually just modify 1 leaf)
- + **Allows search key-sorted scan** of data via linked leaves
- May use **more seeks** vs. other options
- May use **more storage** vs. other options
- **Expensive to combine queries** over multiple B+ trees (e.g. $X > 10$ AND $Y < 20$), as record IDs are unsorted
- **Good for:** high-cardinality attributes, value-sorted record listing, frequent updates

The Hash Index

- Essentially a **hash table on disk** (or other slow medium)
- Key/record pairs stored in **buckets** (one per disk block)
 - Overflow buckets via linked list as needed



Inserting a Record into a Hash Index

- Hash record to **find the bucket**, insert key/record pair
 - Add new overflow bucket if necessary
 - **Overflow hurts performance**
- Use “**dynamic hashing**” to grow hash table as needed
- Example: “**extensible hashing**”
 - In-memory bucket pointer array, double in size as needed
- Example: “**linear hashing**”
 - Add single new bucket when avg. records/bucket is high

Hash Index: Pros and Cons

- + ~1 seek per query (if overflow is limited), better than B+ tree
- + Fewer key compares than B+ tree, usually
- + Inserts/deletes inexpensive (usually just modify 1 bucket)
- Only equality queries supported
- May use more storage vs. other options
- **Good for:** faster equality queries, expensive key compares (e.g., long strings), frequent updates

Summary: Traditional DBMS Indexes

- **B+ Trees**
 - Range/equality queries
 - Inexpensive updates
 - Value-sorted record access
 - Storage space and seeks/query not the best
- **Hash Indexes**
 - Fast equality queries, but no range queries
 - Inexpensive updates
 - Better on seeks, still not best storage space

Overview

- “Traditional” DBMS indexes
- **Bitmap indexes**
- Inverted indexes

Bitmap Indexes

- Less common than B+ trees in DBMSs, but popular in certain applications, e.g., indexing scientific data
- **Idea:** assume r records and k attribute values a_1, \dots, a_k
 - Build k bitmaps b_1, \dots, b_k of length r bits each
 - Bit i in b_j is set iff record i has attribute value a_j

$(RID, value)$

(1, X)

(2, Y)

(3, X)

(4, Z)



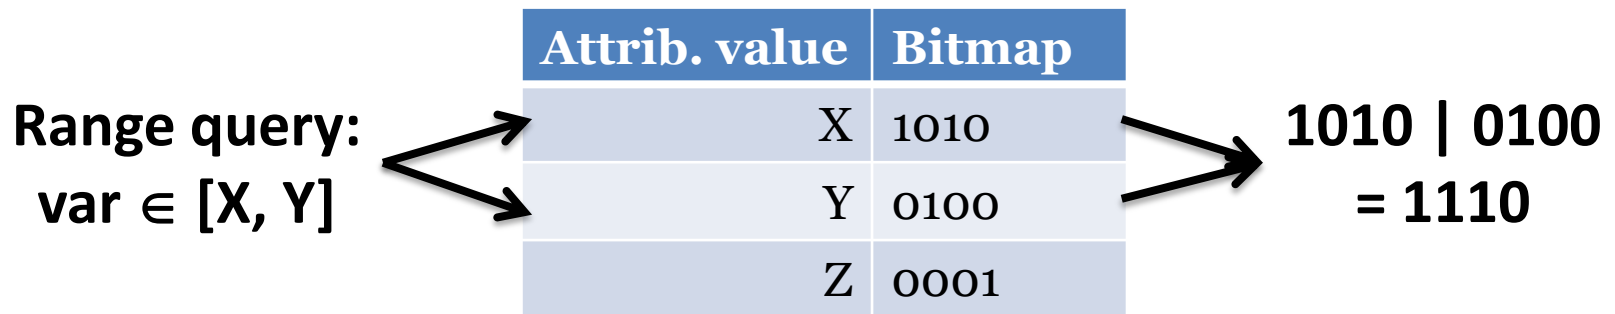
| Attrib. value | Bitmap |
|---------------|--------|
| X | 1010 |
| Y | 0100 |
| Z | 0001 |

Benefits of Bitmap Indexing

- B+ trees/hash indexes allow query on single attributes*
- **Multi-dimensional indexes** exist (R tree, etc.), but quickly succumb to the “**curse of dimensionality**”
 - Combinatorial explosion...
- **Bitmap indexes** are built on **single attributes**, but can be harnessed to answer **multi-attribute queries**
- Also, bitmap indexes may use **less storage** than B+ trees

Querying a Bitmap Index (One Variable)

1. Read the **bitmaps** for values **matching** the query
2. Compute the **bitwise OR** of these bitmaps
3. (Optionally) **convert final bitmap to RIDs**



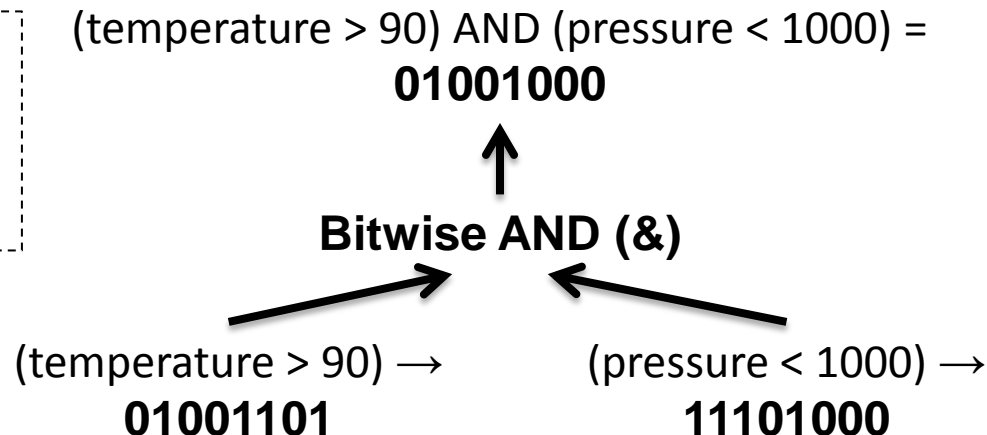
Dataset (RID, value): (1, X), (2, Y), (3, X), (4, Z)

Querying a Bitmap Index (Multivariate)

Example: (temperature > 90) AND (pressure < 1000)

1. Evaluate single-variable constraints to bitmaps, as last slide
2. Combine bitmaps via bitwise operations via expression tree
3. (Optionally) convert final bitmap to RIDs

Dataset (RID, temp., pres.):
(1, 85, 994), (2, 95, 995),
(3, 87, 991), (4, 83, 1002),
(5, 92, 997), (6, 93, 1004),
(7, 81, 1010), (8, 97, 1003)



Updating a Bitmap Index

- **Naïve inserts/deletes** in bitmap indexes would be **costly**
 - Inserting/deleting a bit in the middle shifts many bits
 - Bitmaps are fixed length, would need to expand/shrink
 - **Bitmap compression** exacerbates the situation
- To mitigate, manage where how record RIDs
 - New records have increasing RIDs, always append a bit
 - Deleted record RIDs are not reused; bit position remains
- Still, bitmaps best with **read-mostly/bulk-append** data

3 Key Techniques in Bitmap Indexing

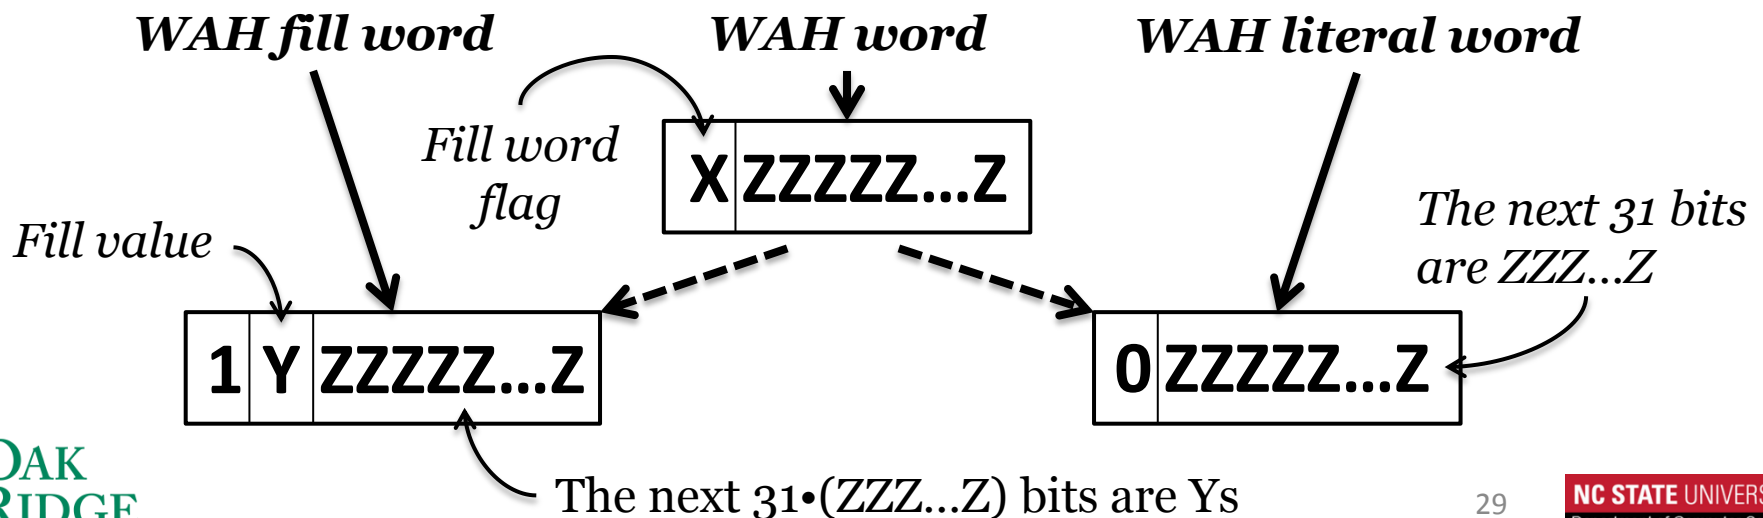
- **Compression**
 - Reduces index size, can speed up query
- **Encoding**
 - Trades index size, query I/O, query CPU time, etc.
- **Binning**
 - Enables indexing of high-cardinality attributes
- These techniques are **orthogonal**, can be **combined**

Compressing Bitmap Indexes

- When an attribute's distinct value count is high (i.e., **high cardinality**), bitmaps become **storage-heavy**
 - However, they also become **sparse** (i.e., mostly 0-bits)
 - Compression can give sublinear growth with cardinality
- Popular compression: **Word-Aligned Hybrid (WAH)**
 - **Run-length encoding** of 0-bit/1-bit runs
 - All operations **aligned to machine words**, for speed
 - AND/OR/NOT possible on WAH without decompression

WAH Bitmap Compression

- Bitmap compressed as a series of 32-bit “**words**”
 - First bit is a “**fill word flag**”; 0→**literal word**, 1→**fill word**
 - **Literal word**: the remaining 31 bits are used verbatim
 - **Fill word**: let f = second bit and the c = the last 30 bits. The fill word represents a $(31 \cdot c)$ -long run of f -bits




Bitmap Index Encodings

- A bitmap index with one bitmap per attribute value is considered “equality encoded”
- Other encoding options exist:
 - **Range encoding**: bit i in b_j set iff record i value $\leq a_j$
 - **Interval encoding**: similar to range, more compact
 - **Binary encoding**: only $\text{ceil}(\log k)$ bitmaps, next slide
 - No longer a 1-1 association between bitmaps and values
- Recover per-attribute bitmaps via bitwise operations

Bitmap Index Encoding Examples


Equality Encoding

Bitmaps (b_1 to b_8)
 (a_k covered \rightarrow RIDs with
 value a_k set in bitmap)




Attribute values (a_1 to a_8)


Range Encoding



Interval Encoding



Binary Encoding



Equality Encoding

(1, X), (2, Y) \rightarrow 00001 (W)
 (3, X), (4, Z) \rightarrow 10100 (X)
 (5, W) \rightarrow 01000 (Y)
 00010 (Z)

Range Encoding

(1, X), (2, Y) 00001
 (3, X), (4, Z) \rightarrow 10101
 (5, W) 11101

Interval Encoding

(1, X), (2, Y) 10101
 (3, X), (4, Z) \rightarrow 10101
 (5, W) 11100

Binary Encoding


(1, X), (2, Y) 01001
 (3, X), (4, Z) \rightarrow 10101
 (5, W)

Why Use Other Bitmap Encodings?

- Each bitmap encoding **optimizes for some property**
- **Range encoding:** **fast range queries**
 - Range query $[a_i, a_j) = b_j - b_i$ (**at most 2 bitmaps needed**)
 - However, index **compression** becomes **less effective**
- **Interval encoding:** **similar to range, fewer bitmaps**
- **Binary encoding:** fewer **bitmaps** = reduced storage
 - However, **every query** requires reading **most bitmaps**

Binning for Bitmap Indexes

- Suppose we want to index a **floating-point attribute**
 - Need one per unique value, which is intractable
- Idea: collect values into “**bins**,” use one bitmap per bin
 - Each “bin” groups together a range of values



Types of Binning

- **Equal-width:** each bin has a fixed width k
 - Easy to compute, sometimes too rigid
- **Equal-frequency:** adjust k bins to each match approx. equal number of records
 - Guarantees balanced binning
 - Requires **sampling** and/or knowledge of **data distribution**
- **Precision:** round values to k significant digits
 - Easy to compute, bin widths adaptive to scale of data
 - May allocate too much precision near 0

The Cost of Binning

- Binning greatly **reduces effective data cardinality**
 - Improves bitmap compression, query performance
- However: **exact query results no longer possible**
 - If a query partially covers a bin, which records match?
- **Option 1:** **Accept error in results**
 - Fast, but (bounded) error in results
- **Option 2:** Do “**candidate checks**” on uncertain records
 - Exact results, but costly random access

Summary: Bitmap Indexing

- Bitmap indexes offer an alternative to traditional DBMS indexes for **efficient multi-attribute querying**
 - Bitmaps may use less storage, thanks to **compression**
 - Bitmap **encoding** can reduce storage, speed up query
 - **Binning** allows bitmap indexing on high-cardinality data
- + Fast equality and range queries + multivariate queries
- + Low storage space
- Frequent inserts/deletes are difficult
- High cardinality data needs binning, with its trade-offs

Overview

- “Traditional” DBMS indexes
- Bitmap indexes
- **Inverted indexes**

Inverted Indexes

- Most commonly used in text/document search engines
- Maps **words** to containing **documents** via a **list of RIDs**
 - This model differs a bit from attributes/records
- Supports “**full text**” search queries: **which documents contain certain words?**

Document 1: “to be or not to be”

Document 2: “to thine own self be true”

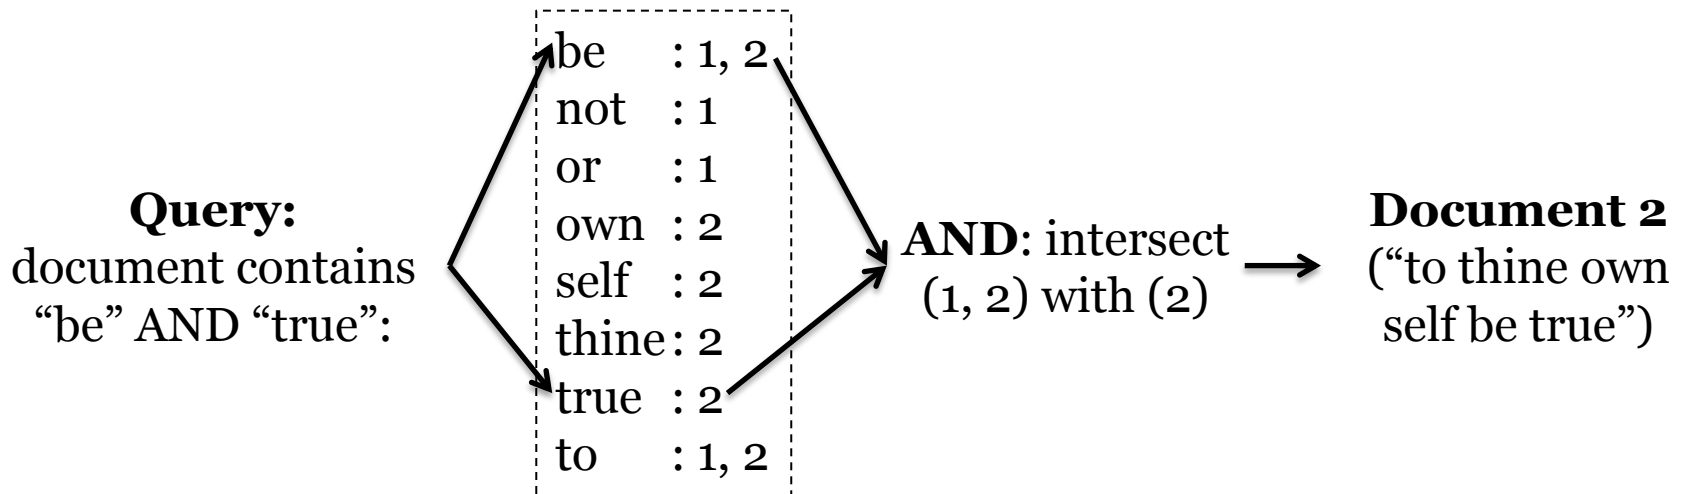
Inverted Index

| | |
|-------|--------|
| be | : 1, 2 |
| not | : 1 |
| or | : 1 |
| own | : 2 |
| self | : 2 |
| thine | : 2 |
| true | : 2 |
| to | : 1, 2 |



Querying an Inverted Index

1. Read the RID lists for each word involved in the query
2. Use list intersections/merges for AND/OR between words



Updating an Inverted Index

- Has some **similar issues to bitmap indexes**
 - RID insert/delete in the middle of a list would shift RIDs
- However, inverted lists needn't be (completely) sorted
 - Inserts can be simply be appended
 - Deleted RIDs may overwritten by moving the last RID
 - Maintaining a partial sort can speed up intersect/merge
- Still, inverted indexes work best for **read-mostly** data

Inverted Index Compression

- Compression greatly reduces storage for long RID lists
- Most compression methods are “**chunk-based**”
 - RIDs grouped into chunks, compressed independently
- **Frame-of-Reference (FOR)**
 - Keep lowest RID, subtract it from all RIDs in chunk
 - Can then use fewer bits per RID

103, 105, 106, 109, 111, 137, 139 →
lowest: 103, relative RIDs: 0, 2, 3, 6, 7, 34, 36 (**6 bits each**)

Inverted Index Compression (cont.)

- **Patched Frame-of-Reference (PFOR)**

- As FOR, but keep **large values** separate as “exceptions”
- If few large exceptions, allows even **fewer bits per RID**

103, 105, 106, 109, 110, 137, 139 →

lowest: 103, relative RIDs: 0, 2, 3, 6, 7, *, * (**3 bits each**), exceptions: 137, 139

- **Patched Frame-of-Reference Delta (PFOR-Delta)**

- As PFOR, but **sort RIDs** and store **pairwise deltas**
- Stores large jumps in the list with only one exception each

103, 105, 106, 109, 110, 137, 139 →

lowest: 103, RID deltas: 2, 1, 3, 1, *, 2 (**2 bits each**), exceptions: 27

Use Case: ElasticSearch*

- Documents are JSON objects with various fields
 - Running example:

```
{  
  content: "The quick brown fox jumped over the lazy dog"  
},  
{  
  content: "Quick brown foxes leap over lazy dogs in summer"  
}
```
- Want to allow **keyword search** within each JSON field
 - Approach: build an **inverted index on each field**
- However, some challenges exist in practice...

Challenge #1: Search Term “Fuzziness”

- Suppose we search for “Quick”
 - Should this match document 1 (below)?
 - Capitalization probably doesn’t matter to the user...
- Other (usually) unimportant differences exist
 - Punctuation
 - Misspellings
 - Synonyms

Document 1: The quick brown fox jumped over the lazy dog

Document 2: Quick brown foxes leap over lazy dogs in summer

Solution: Analysis/Normalization

- Elasticsearch uses “**analysis**” to **normalize** text
 - Applied to **both documents** and **search text** equally
- 1. **Character filters**: cleans up raw text
 - Remove uninteresting punctuation, etc.
- 2. **Tokenizer**: splits text into “**terms**”
 - Split on whitespace, handle hyphenated words, etc.
- 3. **Token filters**: **normalize** each term
 - Example: convert to lowercase
 - Example: remove suffixes (e.g., -ing, -ed)
 - Example: **remove stop word** terms entirely
 - Example: **insert synonym** terms

Challenge #2: Relevance vs. Exact Match

- Suppose we search for “quick brown fox in summer”
 - Should this match document 1 (below)?
 - Maybe, though it is less **relevant** than document 2
- Strict “expression tree” evaluation won’t “**almost match**”

Document 1: The quick brown fox jumped over the lazy dog

Document 2: Quick brown foxes leap over lazy dogs in summer

Solution: Results Ranking

- Elasticsearch “match” queries return documents containing **any query words**, sorted by **relevance score**
 - We can already get the document list: use an OR query
 - Score from matching terms via weighted heuristics:
 1. **Term frequency**: how many **appearances**?
 - More appearances in a field = more weight for matching
 2. **Document frequency**: how **common** is the term?
 - More common across all documents = less weight
 3. **Field-length norm**: how **long** is the **matched field**?
 - Shorter fields (e.g., “title”) = more weight



Other Uses for Inverted Indexes

- Though designed for word-to-document mapping, inverted indexes can still be useful on **structured data**
- Inverted lists have been substituted for bitmap indexes
 - “Words” = attribute values, “documents” = record IDs
- Inverted indexes can use **less storage** than bitmaps
 - Especially high-cardinality attributes; similar to “words”
- Bitmaps are often faster at ANDs/ORs, though

Summary: Inverted Indexes

- Most commonly used to support full text search
- Compression is often used to reduce storage footprint
- Search needs practical considerations beyond indexing
 - + Fast equality queries
 - + Very low storage space for text/high-cardinality data
 - Frequent inserts/deletes are difficult
 - Multivariate query possible, costlier than with bitmaps

Thank You!



Questions?

| Attrib. value | Bitmap |
|---------------|--------|
| X | 1010 |
| Y | 0100 |
| Z | 0001 |

| | |
|-------|--------|
| be | : 1, 2 |
| not | : 1 |
| or | : 1 |
| own | : 2 |
| self | : 2 |
| thine | : 2 |
| true | : 2 |
| to | : 1, 2 |

References

- Much of this presentation's content is covered in:
 “Database Systems: The Complete Book” by Garcia-Molina, Ullman, and Widom
- Classic B-tree paper:
 Douglas Comer. “Ubiquitous B-Tree.” *ACM Comput. Surv.* Vol. 11. No. 2, 1979.
- Bitmap indexing:
 Wu, Kesheng. “FastBit: an efficient indexing technology for accelerating data-intensive science.” *Journal of Physics: Conference Series*. Vol. 16. No. 1, 2005.
 Wu, Kesheng, Ekow J. Otoo, and Arie Shoshani. “Optimizing bitmap indices with efficient compression” *ACM Trans. Database Syst.* Vol. 31, No. 1, 2006.
 Chee-Yong Chan and Yannis E. Ioannidis. “An efficient bitmap encoding scheme for selection queries.” In *Proc. ACM SIGMOD*, 1999.

References (cont.)

- Inverted indexing:

Zobel, Justin, and Alistair Moffat. “Inverted files for text search engines.” *ACM Computing Surveys (CSUR)*. Vol. 38, No. 2, 2006.

Zukowski, Marcin, et al. “Super-scalar RAM-CPU cache compression.” *Data Engineering, 2006. ICDE'06. Proceedings of the 22nd International Conference on*. IEEE, 2006.

Jenkins, John et al. “ALACRITY: Analytics-driven lossless data compression for rapid in-situ indexing, storing, and querying.” *Transactions on Large-Scale Data-and Knowledge-Centered Systems X*, 2013.

- Elasticsearch-related content can be found at:

<http://www.elastic.co/guide/en/elasticsearch/guide>